
Cooperation in Multi-Agent Learning

Karush Suri, Lacro Pavel

Department of Electrical & Computer Engineering, University of Toronto, Canada.
karush.suri@mail.utoronto.ca

Abstract

Advancements in Multi-Agent Reinforcement Learning (MARL) are motivated by cooperation in agents arising from Game Theory (GT). Agents must collaborate in practical scenarios in order to achieve complex objectives and attain strategies which depict optimal behavior. The need for cooperation is further highlighted in the case of partially-observed settings wherein agents have restricted access to environment observations. We revisit cooperation in MARL from the viewpoint of GT and stochastic dynamics of environments. The contributions of our work are twofold. (1) We analyze and demonstrate the effectiveness of cooperative MARL in the case of complex and partially-observed tasks consisting of high-dimensional action spaces and stochastic dynamics. (2) We leverage the empirical demonstrations to construct a novel optimization objective which addresses the detrimental effects of spurious states across agents. Our large-scale experiments carried out on the StarCraft II benchmark depict the effectiveness of cooperative MARL and our novel objective for obtaining optimal strategies under stochastic dynamics.

1 Introduction

2 Related Work

2.1 Learning in Games

2.2 Multi-Agent Learning

3 Preliminaries

3.1 Stochastic Markov Games

3.2 Q-Learning

3.3 Multi-Agent Learning

4 Cooperation in Multi-Agent Learning

4.1 The Partial Observability Setting

4.2 Learning Model-Free Behaviors

5 Tackling Spurious Dynamics

6 Experiments

6.1 The StarCraft II Benchmark

6.2 Performance

6.3 Spurious Dynamics

7 Conclusion

References